QALA

Universal Solution Management Operating System

Request for Comments (RFC) | v2.0 | Universal Edition

Field	Value
Document Type	RFC — Universal Edition v2.0
System	Qala Universal Solution Management OS
Status	Open for Comment — 30-day comment period
RFC Count	22 RFCs across architecture, solution model, domain extensions, and strategy
Purpose	Solicit structured feedback on design decisions for the Universal Solution Management OS covering all solution types and industries
Audience	Engineering, Architecture, Domain Experts, Security, Product, External Reviewers

1 Introduction and How to Use This RFC

This document is a formal Request for Comments on the design of the Qala Universal Solution Management Operating System. It presents 22 open design questions spanning platform architecture, the universal solution model, domain-specific extensions, integration, compliance, AI, and open source strategy. Feedback shapes the platform before decisions are locked.

An RFC is a structured proposal that explicitly names its open questions and alternative approaches. The goal is not to present a finalized design but to invite rigorous challenge, domain expertise, and creative alternatives from the broadest possible community of reviewers.

> *How to respond: Submit comments via the project issue tracker or document annotations. Reference the RFC ID and Question ID (e.g. Q-ARCH-001) in every comment. The Architecture Review Board will respond to every substantive comment within the 30-day period.*

1.1 RFC Index — All 22 RFCs

RFC ID	Title	Domain	Status
RFC-001	Polyglot Architecture and Language Selection	Architecture	Open
RFC-002	SDE Model: Universal Environment Types	Platform	Open
RFC-003	Hermetic Build and Attestation for All Solution Types	Build / Security	Open
RFC-004	Universal Event Architecture and Kafka Topology	Platform	Open
RFC-005	AI Agent: Universal Intelligence Architecture	AI / Intelligence	Open
RFC-006	Universal Solution Type System and Extensibility	Solution Model	Open
RFC-007	Change Control Governance: Universal CCR Model	Governance	Open
RFC-008	Universal Quality Management Model	Quality	Open
RFC-009	Universal Release and Distribution Management	Release	Open
RFC-010	Physical Goods, CPG and Manufacturing Domain Extension	Domain Extension	Open
RFC-011	Financial, Investment and Capital Solution Domain Extension	Domain Extension	Open
RFC-012	Agricultural and Agribusiness Solution Domain Extension	Domain Extension	Open
RFC-013	Tools, Toolchains, Libraries, Packages and SDK Domain Extension	Domain Extension	Open
RFC-014	Service, Consulting and Professional Services Domain Extension	Domain Extension	Open

RFC-015	Business, Process and Operational Solution Domain Extension	Domain Extension	Open
RFC-016	Research, Academic and Scientific Solution Domain Extension	Domain Extension	Open
RFC-017	Legal, Tax and Regulatory Solution Domain Extension	Domain Extension	Open
RFC-018	Universal Compliance Framework Architecture	Compliance	Open
RFC-019	Physical World Integration: ERP, MES, LIMS, IoT Bridges	Integration	Open
RFC-020	Multi-Tenancy, Tenant Isolation and Cross-Tenant Sharing	Platform	Open
RFC-021	Identity, RBAC and Electronic Signature Architecture	Security	Open
RFC-022	Open Source, Ecosystem and Commercialisation Strategy	Strategy	Open

2 Architecture Context for Review

The expanded architecture adds a sixth Domain Extension Plane to support universal solution types without polluting the platform core. All domain-specific logic lives in independently deployable, community-contributable Domain Packs.

Plane	Language	Responsibilities
Kernel Plane	Rust	Platform integrity, service registry, health monitoring, event aggregation, cross-domain solution graph engine, platform self-governance.
Control Plane	Go	SDE Management, Solution Factory, Identity, Notifications, Change Control (CCR/CCB), Bug Tracking, Release Management, Universal Solution Lifecycle State Machine.
Execution Plane	Go / Scala	CI/CD Engine, Test Management, Benchmarking, Hermetic Build, Artefact Registry, Distribution Engine. Go for orchestration; Scala/Akka for reactive streaming where needed.
Intelligence Plane	Rust	Universal AI Agent (domain-aware), SEM, Vulnerability Management, Privacy Management. Domain-specific intelligence models hot-swappable per solution type.
Platform Operations	Go / Terraform	Environment Management, Tool Registry, Vendor Integration, Backup and Recovery, Audit and Compliance Reporting.
Domain Extension Plane	Go + YAML DSL	Vertical-specific solution type schemas, compliance framework adapters, physical world integration bridges, domain workflow templates, industry vocabulary overlays, role-specific UX personas.

> *Domain Packs are the primary extension mechanism. A Domain Pack for CPG defines: the Good/Product/ProductLine solution type schemas; the formulation CCR workflow template; GMP and REACH compliance evidence mappings; ERP integration bridge configuration; and the CPG-specific portal vocabulary. Domain Packs are versioned, independently deployable, and community-contributable.*

3 RFC-001: Polyglot Architecture and Language Selection

RFC-001 — Polyglot Architecture and Language Selection
Status: Open for Comment **Domain:** Architecture

The six-plane architecture uses four languages: Rust (Kernel + Intelligence), Go (Control + Platform Ops + Domain Ext), Scala (Execution streaming), and YAML DSL (domain pack schemas). This RFC asks whether the polyglot complexity is justified by the performance and safety benefits.

Option	Languages	Pros	Cons
Proposed: Polyglot	Rust + Go + Scala + YAML	Best tool per domain. Rust safety for security-critical code. Scala for reactive pipelines.	Four languages. High hiring bar. Complex CI/CD across language ecosystems.
Alt A: Go-only	Go + YAML	Single language. Fastest iteration. Largest talent pool. Simplest toolchain.	Go lacks Rust's safety guarantees. Less natural for reactive streaming vs. Scala.
Alt B: Rust + Go	Rust + Go + YAML	Best of two worlds. Go productivity for business logic; Rust safety for critical paths.	Still polyglot. Rust learning curve. Slower initial velocity.
Alt C: JVM-first	Kotlin + Java + YAML	Strong JVM ecosystem. Kotlin is modern, safe, concise. Excellent library support.	JVM startup time for Kernel/SEM latency. GC pauses in security-critical code.

Q-ID	Open Question — Requires Decision	Decision By
Q-ARCH-001	Is the reactive streaming capability of Scala/Akka worth the language diversity cost? Could Go + Temporal (workflow orchestration) deliver equivalent pipeline orchestration capability with a single-language stack?	Pre-alpha build
Q-ARCH-002	For the Kernel Plane, is Rust's memory safety genuinely necessary given that Go's memory safety is "good enough" and Go's garbage collector has improved substantially in recent versions?	Pre-alpha build
Q-ARCH-003	Should the Domain Extension YAML DSL be expressive enough for complex domain-specific validation logic, or should Domain Pack authors write Go code? A pure YAML DSL limits expressiveness; requiring Go raises the contributor bar.	Pre-alpha build
Q-ARCH-004	What is the minimum viable language set that preserves the key architectural benefits (safety for security code, streaming for pipelines, speed for control plane) while minimising polyglot overhead?	Pre-alpha build

4 RFC-002: SDE Model — Universal Environment Types

RFC-002 — SDE Model: Universal Environment Types
Status: Open for Comment **Domain:** Platform

The Solution Development Environment (SDE) must serve as the governed workspace for creating any solution type. For software, this is a hermetic build container. For a pharmaceutical formulation team, it is a simulation and regulatory tooling workspace. For a quantitative finance team, it is a backtesting and data access workbench. For a farm operations team, it is an IoT firmware build and sensor configuration environment.

SDE Type	Proposed Implementation	Key Requirement	Primary Challenge
Software Dev	OCI container + Kubernetes orchestration	Hermetic builds, full IDE integration, fast provisioning.	None — well-understood, mature space.
Lab / Formulation	VM or Kata container + NFS mounts	Access to commercial simulation software, file-based toolchains.	Commercial software licensing in containerised environments.
Financial Research	OCI container + optional GPU attachment	Large dataset access, quantitative libraries (Python, R, Julia).	GPU attachment; data access latency for large datasets.
Edge / IoT / AgriTech	Nix flake + bare-metal / microVM	Offline operation, cross-compilation, constrained resources.	Nix maturity; air-gap dependency pre-population strategy.
Knowledge / Service	Lightweight OCI container + CMS integration	Document versioning, AI retrieval, knowledge base access.	CMS integration model; knowledge graph consistency.
Research / Academic	OCI + large storage mounts + dataset registry	Dataset versioning, experiment reproducibility, compute scaling.	Petabyte-scale dataset management; HPC integration.
Business / Process	Lightweight OCI + process modelling tools	BPMN/flowchart tooling, document management, collaboration.	Process model storage and versioning format selection.

Q-ID	Open Question — Requires Decision	Decision By
Q-SDE-001	Should "SDE" be a single unified concept with modular capability packs (hermetic-build-pack, dataset-access-pack, knowledge-workspace-pack) or should it be subtyped explicitly (SoftwareSDE, LabSDE, ResearchSDE, KnowledgeSDE)? The unified model is simpler to govern; subtypes are more intuitive for non-software users.	Pre-alpha
Q-SDE-002	For Lab and Research SDEs accessing commercial licensed software (Matlab, ANSYS, SAS), how should license compliance be managed in containerised environments? Node-locked and	Pre-beta

	seat-licensed software fundamentally conflicts with ephemeral container models.	
Q-SDE-003	What does an SDE "snapshot" mean for a non-software environment? For a software SDE, a snapshot captures filesystem + config + dependencies. For a formulation lab SDE, what constitutes the complete capturable state?	Pre-alpha build
Q-SDE-004	Should non-software SDE types be in scope for v1, or should v1 ship with Software SDE only, with domain SDE types added via Domain Packs in subsequent releases? Phasing reduces v1 scope risk but delays non-software adoption.	Pre-alpha

5 RFC-003: Hermetic Build and Attestation for All Solution Types

RFC-003 — Hermetic Build and Attestation for All Solution Types
Status: Open for Comment **Domain:** Build / Security

SLSA (Supply-chain Levels for Software Artefacts) provides a robust build attestation framework for software. For the universal platform, the equivalent operation across all solution types must be defined. A "build" for a pharmaceutical formulation is a simulation run and regulatory pre-check. For a financial model, it is a calibration and backtesting run. For a research experiment, it is a pipeline execution. Each requires an attestation that the result was produced from known, approved inputs in a governed environment.

Solution Type	Build Equivalent	Reproducibility Class	Proposed Attestation Record
Software / Firmware	Compile, link, package in hermetic container.	Deterministic — bit-for-bit reproducible.	Source commit, env hash, dependency digests, output hash, signer identity.
Physical Formulation	Simulation run + specification validation + regulatory pre-check.	Deterministic — same inputs produce same simulation outputs.	Formulation version, simulation tool version, input parameters, output report hash.
Financial Model	Calibration run + backtesting + parameter validation.	Partially deterministic — market data inputs can vary.	Model version, data snapshot ID, calibration parameters, backtest report hash, risk metric outputs.
ML / AI Model	Training run + evaluation + serialisation.	Non-deterministic — stochastic gradient descent.	Dataset version, hyperparameters, random seed, evaluation metrics, model artefact hash.
Research Experiment	Pipeline execution + data processing + result generation.	Deterministic if environment is hermetic.	Script version, dataset version, env hash, result hashes, execution log hash.
Tax Rule Set	Rule compilation + validation suite + jurisdiction test run.	Deterministic.	Rule set version, test suite version, jurisdiction variants, validation report hash.
Service Methodology	Document assembly + review workflow completion + approval.	Deterministic — document content is fixed at approval.	Document version, approver identities, approval timestamps, document hash.
Business Solution	Process model validation + stakeholder review + acceptance.	Deterministic at the process model level.	Process model version, validation results, acceptance records, model hash.

Q-ID	Open Question — Requires Decision	Decision By

Q-HB-001	Should Qala publish a "Universal Solution Attestation Schema" as an open standard extending SLSA to cover all non-software solution types? This would be a significant community contribution and could drive cross-industry adoption, but requires substantial standards body engagement.	Post-alpha
Q-HB-002	For stochastic processes (ML training, financial model calibration), bit-for-bit reproducibility is not achievable. Should the attestation standard require "governed reproducibility" (same environment + same inputs produces results within defined tolerance bounds) rather than exact reproducibility?	Pre-beta
Q-HB-003	Should SLSA Level 3 compliance be mandatory for all solution types from launch, or should the hermetic build requirement be phased — mandatory for software from v1, voluntary for other solution types until domain-specific attestation standards are defined?	Pre-alpha build
Q-HB-004	For physical goods, the "build" is a physical manufacturing process that Qala cannot govern directly. How should attestation connect the digital specification version to the physical batch record? Should Qala define a "specification-to-batch linkage" attestation that connects the digital record to the physical production event?	Pre-GA

6 RFC-004: Universal Event Architecture

RFC-004 — Universal Event Architecture and Kafka Topology
Status: Open for Comment **Domain:** Platform

Topic	Domain	Retention	Key Event Types
SDE_EVENTS	Platform	30 days	SDE_CREATED, SNAPSHOTTED, QUARANTINED, ROLLED_BACK, ARCHIVED, DECOMMISSIONED
SOLUTION_EVENTS	Universal	90 days	SOLUTION_CREATED, VERSION_RELEASED, STATE_CHANGED, RECALLED, RETIRED
BUILD_EVENTS	Execution	90 days	BUILD_STARTED, PASSED, FAILED, ATTESTATION_CREATED, DEPENDENCY_VIOLATION
RELEASE_EVENTS	Execution	90 days	RELEASE_PLANNED, GATE_PASSED, GATE_FAILED, DEPLOYED, ROLLED_BACK, RECALLED
QUALITY_EVENTS	Universal	1 year	TEST_PASSED, TEST_FAILED, DEFECT_CREATED, QUALITY_GATE_PASSED, QUALITY_GATE_FAILED
CHANGE_EVENTS	Universal	1 year	CCR_SUBMITTED, CCR_APPROVED, CCR_REJECTED, CHANGE_IMPLEMENTED, DRIFT_DETECTED
SECURITY_EVENTS	Security	1 year	THREAT_DETECTED, QUARANTINED, VULNERABILITY_FOUND, SUPPLY_CHAIN_VIOLATION
COMPLIANCE_EVENTS	Compliance	7 years	POLICY_CHECKED, VIOLATION_DETECTED, EVIDENCE_GENERATED, REPORT_PUBLISHED
DISTRIBUTION_EVENTS	Distribution	1 year	DISTRIBUTED, CHANNEL_UPDATED, DISTRIBUTION_SUSPENDED, RECALL_INITIATED
AI_EVENTS	Intelligence	14 days	RECOMMENDATION_PUBLISHED, PREDICTION_ISSUED, ANOMALY_DETECTED, FEEDBACK_RECORDED
DOMAIN_EVENTS	Domain Ext.	1 year	Domain-specific: BATCH_RELEASED, FORMULATION_APPROVED, FUND_INCEPTED, HARVEST_RECORDED
AUDIT_EVENTS	Compliance	7 years	Every state-modifying action — append-only, cryptographically chained

Q-ID	Open Question — Requires Decision	Decision By
Q-EVT-001	AUDIT_EVENTS requires 7-year immutable retention — this exceeds Kafka's practical strengths. Should audit events write to a dedicated append-only, cryptographically hash-chained audit log store (separate from Kafka) with Kafka as the real-time distribution bus only?	Pre-alpha build
Q-EVT-002	DOMAIN_EVENTS is currently proposed as a single shared topic with domain-type discriminators. Should each vertical instead have its own topic (CPG_EVENTS, AGRI_EVENTS,	Pre-alpha build

	FINANCIAL_EVENTS) for consumer filtering efficiency? What is the topic proliferation trade-off at scale?	
Q-EVT-003	How should schema governance work for DOMAIN_EVENTS that are defined by Domain Pack authors? Should every Domain Pack submit an AsyncAPI spec that is validated against the platform Schema Registry before the pack is activated in any SDE?	Pre-beta
Q-EVT-004	Physical goods supply chain events (BATCH_RECEIVED, COLD_CHAIN_VIOLATION, PRODUCT_DISPATCHED) may arrive from external systems (ERP, MES, logistics platforms) rather than from Qala services. How should external event ingestion be authenticated, validated, and integrated into the event bus without contaminating the internal event trust model?	Pre-beta

7 RFC-005: AI Agent — Universal Intelligence Architecture

RFC-005 — AI Agent: Universal Intelligence Architecture
Status: Open for Comment **Domain:** AI / Intelligence

The Qala AI Agent is not a general-purpose language model assistant. It is a lifecycle intelligence engine with full context about every solution it manages — the solution's history, quality trajectory, security posture, relationship graph, domain rules, and current state. Its capabilities are domain-aware: it generates software tests, predicts pharmaceutical batch quality failures, detects financial model drift, recommends agronomic interventions, retrieves consulting methodologies, and optimises research experiment designs — all from one architecture.

Domain	AI Capability	Input Data	Output Type
Software	Test generation, defect prediction, pipeline optimisation, anomaly detection, security risk scoring.	Build history, test results, defect records, code changes, benchmarks.	Test cases, risk scores, pipeline recommendations, anomaly alerts.
Physical Goods	Batch quality prediction, formulation optimisation, recall risk detection, yield optimisation.	Batch records, quality tests, recall history, supplier records, sensor data.	Quality risk alerts, formulation suggestions, recall probability scores.
Financial	Model drift detection, compliance anomaly detection, pricing model validation, portfolio risk scoring.	Model performance history, regulatory findings, market data, backtest results.	Model risk alerts, compliance predictions, pricing validation reports.
Agricultural	Yield prediction, crop disease detection, irrigation optimisation, supply chain risk scoring.	Sensor data, weather forecasts, historical yields, satellite imagery, logistics data.	Yield forecasts, field intervention recommendations, logistics optimisation.
Professional Svcs	Methodology recommendation, outcome prediction, engagement quality scoring, knowledge retrieval.	Engagement records, outcomes, client NPS, knowledge base, methodology versions.	Methodology matches, outcome forecasts, knowledge summaries, quality alerts.
Research	Experiment outcome prediction, reproducibility verification, anomaly detection in results.	Publication records, experiment results, datasets, simulation runs.	Reproducibility risk alerts, experiment design recommendations.
Legal / Tax	Regulatory change impact prediction, compliance gap detection, jurisdiction variant recommendation.	Regulatory databases, product version history, examination findings, legislative feeds.	Impact assessments, compliance gap reports, update prioritisation.
Business / Ops	Process performance prediction, bottleneck	Process telemetry, outcome metrics, change history, benchmark data.	Process improvement recommendations, change

	identification, change risk scoring.		risk scores, outcome forecasts.

Q-ID	Open Question — Requires Decision	Decision By
Q-AI-001	Should the AI Agent use a single universal model fine-tuned across all domain data, or domain-specific models (SoftwareAI, CPG_AI, FinancialAI) each trained on domain-specific data? Universal model benefits from cross-domain learning; domain models provide higher quality per vertical but require separate training pipelines.	Pre-beta
Q-AI-002	For regulated domains (pharma, finance, medical devices), AI recommendations that influence product decisions may be subject to explainability requirements (EU AI Act, FDA AI/ML guidance). How should the AI Agent produce explainable outputs for regulated solution types while preserving model performance?	Pre-GA
Q-AI-003	How should the AI Agent handle domains where it has insufficient training data at launch (e.g., a new vertical)? Should it decline to make recommendations, surface low-confidence recommendations with explicit uncertainty quantification, or fall back to rule-based heuristics?	Pre-beta
Q-AI-004	Should Qala's domain intelligence be exposed as an API so customers can integrate Qala's solution lifecycle intelligence into their own products (e.g., a CPG company integrating Qala's batch quality prediction into their production control system)?	Post-GA
Q-AI-005	What is the data governance model for AI training data derived from customer solution records? Who owns the trained model improvements — Qala or the customer? How should model improvement be apportioned when it is driven by cross-customer learning?	Pre-GA

8 RFC-006: Universal Solution Type System

RFC-006 — Universal Solution Type System and Extensibility
Status: Open for Comment **Domain:** Solution Model

The Universal Solution Type System is the foundational data model of Qala. Every solution managed on the platform is an instance of a solution type. Every solution type inherits universal attributes and adds domain-specific attributes. The type system must be coherent enough to govern universally and flexible enough to represent every identified solution category.

8.1 The Type Hierarchy

Level	Description	Examples
Solution Category (10)	Top-level grouping of related solution types.	Software, Physical Goods, Services, Financial, Agricultural, Business, Research, Resources, Platforms, Legal/Tax
Solution Type (40+)	First-class managed entity type with type-specific lifecycle state machine.	Application, Good, Financial Instrument, Agricultural Platform, Consulting Service, Research Asset, Tax Solution
Solution Subtype (extensible)	Domain-specific specialisation of a type, added via Domain Pack.	ConsumerPackagedGood (subtype of Good), ETF (subtype of Financial Instrument), CropManagementSystem (subtype of AgriculturalPlatform)

8.2 Universal Solution Attributes

Attribute Group	Attributes	All Types?
Identity	ID, name, type, subtype, owner, organisation, created_at, created_by.	Yes
Versioning	semantic_version, version_history, changelog, version_status.	Yes
Lifecycle	lifecycle_state (per type FSM), state_history, last_transition, transition_reason.	Yes
Provenance	creation_context, change_history (linked CCR IDs), distribution_history.	Yes
Artefacts	artefact_manifest (all digital artefacts, versioned, with attestation references).	Yes
Quality	quality_record_refs, quality_gate_history, benchmark_history.	Yes
Relationships	dependencies, components, parent_solutions, child_solutions, cross-type relationships.	Yes
Compliance	applicable_frameworks, compliance_status_per_framework, evidence_refs.	Yes

Domain Metadata	Type-specific attributes defined by Domain Pack schema.	Type-specific

Q-ID	Open Question — Requires Decision	Decision By
Q-MODEL-001	The taxonomy currently has 10 categories and 40+ types. Should types be a strict flat list (simpler to reason about) or a hierarchy (categories > types > subtypes)? A flat list avoids hierarchy complexity; a hierarchy enables incremental onboarding (learn Application before learning all 40+ types).	Pre-alpha
Q-MODEL-002	Should solution type subtypes be user-definable from launch, or should the initial release have a fixed, curated set of types and subtypes, with user-defined extension deferred to v2? Fixed types reduce onboarding complexity but may not fit every domain use case.	Pre-alpha
Q-MODEL-003	Quality is a universal attribute, but "quality" means fundamentally different things per domain. Should the quality record be a generic key-value store (domain-defined quality metrics), or should Qala define a typed quality model per solution category?	Pre-alpha build
Q-MODEL-004	Cross-type solution relationships (e.g. a Platform composed of Applications and Services that produce Physical Goods) are central to the universal model. Should these relationships be a typed directed graph (most powerful), a simple parent-child hierarchy (simpler), or a flat tag-based association (most flexible but least governed)?	Pre-alpha build
Q-MODEL-005	Should "Product Line" and "Solution Family" be first-class solution types (managed entities with their own lifecycle), or should they be modelled as relationship groupings of individual product solutions? A first-class type has its own CCR workflow; a grouping is simpler but may miss governance needs.	Pre-beta

9 RFC-007: Change Control Governance — Universal CCR Model

RFC-007 — Change Control Governance: Universal CCR Model
Status: Open for Comment **Domain:** Governance

Every change to every solution — a code commit, an ingredient substitution in a formulation, a financial model parameter recalibration, an SLA term update, a tax rule modification, a process redesign — must flow through the Change Control Request (CCR) workflow. The CCR model must be universal enough to govern all change types, flexible enough to reflect the very different risk profiles of a routine dependency upgrade vs. a pharmaceutical ingredient substitution.

Change Type	Risk Profile	Proposed Approval Path	Regulatory Requirement
Routine software dependency upgrade (patch)	Low	Single senior engineer approval + automated validation.	None
Software API breaking change	Medium	Tech Lead + Release Manager + automated consumer impact check.	Varies by contract
CPG formulation ingredient substitution	High	Product Dev + QA + Regulatory Affairs + mandatory validation run.	FDA/EFSA change notification or re-registration
Financial model parameter recalibration	High	Quant Lead + Risk Manager + Model Validation Committee.	MiFID II model governance requirements
Tax rule set update (new legislation)	High	Tax Product Manager + Legal + Compliance + effective date scheduling.	Must match legislative effective date exactly
Agricultural IoT firmware update	Medium	Engineering Lead + QA + field impact assessment.	May require regulatory notification for certain device classes
Service SLA term change	Medium	Service Architect + Legal + Customer Success.	Contract-specific; customer notification required
Research pipeline algorithm update	Low-Medium	Principal Investigator + peer review gate.	Publication reproducibility requirements
Business process redesign	High	Process Owner + Impacted stakeholders + Change Champion.	Varies by regulated industry context

Q-ID	Open Question — Requires Decision	Decision By
Q-CM-001	Should the CCR risk scoring model be universal (same algorithm for all solution types) or domain-specific (a pharmaceutical CCR has different inherent risk factors than a software CCR)? A	Pre-beta

	universal model is simpler to govern; a domain-specific model better reflects actual risk.	
Q-CM-002	For regulated solution types (pharma, financial, tax), some change types carry regulatory obligations (FDA notification, MiFID II governance procedures) that must be embedded in the CCR workflow. Should these regulatory obligations be hard-coded into Domain Pack CCR templates, or should they be configurable per organisation based on their specific regulatory exposure?	Pre-beta
Q-CM-003	Should the CCR workflow support multi-stage sequential approval (e.g. Scientific Committee > Regulatory Affairs > Legal > Executive sign-off for a pharmaceutical product change) with each stage having its own evidence requirements and approval criteria?	Pre-beta
Q-CM-004	How should emergency changes be governed for non-software solution types? A software emergency change can be rolled back in minutes. A pharmaceutical emergency change (product recall) or financial emergency change (instrument repricing) may have irreversible consequences. Should the emergency change process have type-specific constraints for irreversible solution types?	Pre-beta

10 RFC-008: Universal Quality Management Model

RFC-008 — Universal Quality Management Model
Status: Open for Comment **Domain:** Quality

Quality management in Qala must work across every solution type. The concepts of "test," "defect," "quality gate," and "benchmark" apply universally — but their specific implementations vary dramatically across domains. A software test suite, a pharmaceutical batch quality test, a financial model backtesting run, a field trial for an agricultural solution, and a peer review for a research publication are all "quality validation" events. Qala must provide a universal quality framework that captures all of these coherently.

Quality Concept	Software	Physical Goods	Financial	Agricultural	Research
Test Suite	Unit + integration + system tests	Batch quality tests, stability studies, regulatory tests	Backtesting runs, stress tests, model validation suite	Field trials, efficacy studies, environmental tests	Replication runs, peer review, statistical validation
Defect	Code defect, failed assertion, security finding	Non-conformance, contamination, specification deviation	Model error, risk limit breach, data error	Crop failure, yield underperformance, disease outbreak	Irreproducible result, data integrity failure, methodology error
Quality Gate	Coverage threshold, SAST pass, no critical bugs	GMP compliance, spec conformance, regulatory pre-check	Risk model validity, backtest threshold, compliance check	Efficacy threshold, safety clearance, regulatory approval	Reproducibility score, statistical significance, peer review pass
Benchmark	Test pass rate, defect density, DORA metrics	Batch quality yield, defect rate, recall frequency	Sharpe ratio, tracking error, VaR accuracy	Yield vs. forecast, intervention efficacy rate	Reproducibility rate, citation impact, data reuse rate

Q-ID	Open Question — Requires Decision	Decision By
Q-QM-001	Should "quality record" be a generic schema (key: metric_name, value: metric_value, type: PASS/FAIL/SCORE) applicable to all solution types, or should each solution category have a typed quality schema with strongly-typed metrics? Generic is universal	Pre-alpha build

	but loses domain-specific validation; typed is safer but harder to extend.	
Q-QM-002	For physical goods, quality tests are conducted in the physical world (lab tests, field trials) and the results must be imported into Qala. Should Qala provide a Quality Event Ingestion API that accepts results from LIMS, lab instruments, and field testing tools via a standard event format?	Pre-beta
Q-QM-003	How should cross-domain quality benchmarking work? Is it meaningful to compare the "quality trajectory" of a software product with the "quality trajectory" of a CPG product? If so, what normalised quality metrics apply universally? If not, should benchmarking be strictly domain-scoped?	Pre-beta

11 RFC-009: Universal Release and Distribution Management

RFC-009 — Universal Release and Distribution Management
Status: Open for Comment **Domain:** Release

Release management governs the promotion of a solution version from development through quality gates to distribution. Distribution management governs the delivery of that released version to its consumers or channels. Both concepts apply universally across all solution types.

Solution Type	Release Gate Types	Distribution Channel	Rollback / Recall Model
Software Application	Test, Coverage, Security, Performance, Attestation, Approval.	App store, Cloud deploy, Package registry.	Instant rollback via Blue/Green; previous version reactivation.
Physical Good (CPG)	Quality test, Regulatory, Stability, Shelf-life, Compliance, Approval.	Retail channels, Distribution centres, DTC.	Recall — batch identification, channel notification, physical retrieval.
Financial Instrument	Risk model, Compliance (MiFID II/SEC), Legal, Management approval.	Fund inception, Bloomberg, Exchange listing.	Product suspension — investor notification, redemption window.
Tax Solution	Rule validation, Jurisdiction test, Legal review, Effective date gate.	Software update channel, API update, Portal publish.	Rule rollback to prior version; jurisdiction override.
Agricultural Solution	Field trial, Regulatory (EPA/EFSA), Safety, Seasonal timing.	Farm deployment, AgriRetailer, Direct to farmer.	Solution rollback, field protocol update, advisory withdrawal.
Consulting Service	Content review, Legal, Partner sign-off, Client pilot approval.	Service catalogue, Sales enablement, Client delivery.	Version revert, engagement protocol update.
Research Asset	Peer review, Data validation, Reproducibility check, IRB approval.	Publication, Data repository, Open source.	Retraction notice, dataset version revert.

Q-ID	Open Question — Requires Decision	Decision By
Q-REL-001	The release gate framework currently defines mandatory gates as platform defaults. For non-software solution types, the mandatory gates will be domain-specific (a CPG release must have a Regulatory Gate; a software release must have a Security Gate). Should mandatory gates be defined by Domain Packs, configurable per solution type, or hardcoded per category?	Pre-beta
Q-REL-002	For physical goods, "rollback" is a physical recall — a fundamentally different operation from a software deployment reversal. Should Qala model Recall as a specialised workflow	Pre-beta

	distinct from software Rollback, with its own event types, notification templates, and regulatory reporting obligations?	
Q-REL-003	Some releases have hard external effective dates (tax rule changes, regulatory compliance deadlines, legislative effective dates). Should the release pipeline support "effective date gates" that hold a release in a deployed-but-inactive state until a scheduled activation timestamp?	Pre-beta
Q-REL-004	For distributed physical goods, the "distribution record" must track which specific batch went to which specific retailer/distributor/consumer. How granular should Qala's distribution tracking be, and how should it integrate with retail supply chain systems (GS1 EPCIS, EDI)?	Pre-GA

12 RFC-010 to RFC-017: Domain Extension RFCs

These RFCs address the design of Domain Packs for each major industry and solution category. Each RFC identifies the key open questions that must be resolved before the Domain Pack can be built.

RFC-010 — Physical Goods, CPG and Manufacturing Domain Extension

Status: Open for Comment **Domain:** Domain Extension

Q-ID	Open Question — Requires Decision	Decision By
Q-CPG-001	Should Qala natively support GS1 barcodes, GTIN identifiers, and GS1 EPCIS event format as the identity and traceability standard for CPG and physical goods solutions, enabling direct integration with retailer supply chain systems without custom integration work?	Pre-beta
Q-CPG-002	Recall management for physical goods involves regulatory notification (FDA MedWatch, EFSA RASFF, country-specific portals), supply chain partner notification, and consumer notification — all with regulatory time windows. Should Qala's recall workflow have built-in regulatory notification connectors, or treat notification as a generic distribution event?	Pre-beta
Q-CPG-003	For pharmaceutical products, 21 CFR Part 11 requires qualified electronic signatures and validated audit trails beyond Qala's standard audit model. Should Qala offer a "Part 11 Mode" that applies enhanced signature, audit, and validation requirements to a specific SDE or Solution Factory?	Pre-GA
Q-CPG-004	Product Line governance requires cascading change impact: changing a shared ingredient affects all products in the line. Should Qala's CCR workflow automatically compute and surface affected solutions when a component-level change is proposed, before the CCR is approved?	Pre-beta

RFC-011 — Financial, Investment and Capital Solution Domain Extension

Status: Open for Comment **Domain:** Domain Extension

Q-ID	Open Question — Requires Decision	Decision By
Q-FIN-001	Financial regulators (SEC, FCA, ESMA) may conduct examinations and request access to governance evidence. Should Qala support a "Regulator Access" RBAC role that grants time-limited, read-only access to a scoped subset of governance and audit records — with the access itself fully audited?	Pre-GA
Q-FIN-002	Investment mandates and financial product terms are legal documents with version histories and amendment schedules. Should Qala's Document Registry (part of the Artefact Registry) support legal document types with redlining, amendment tracking, and electronic signature integration?	Pre-beta

Q-FIN-003	Trading algorithms require a "paper trading" (simulation) stage before live deployment, analogous to a staging environment for software. Should Qala's deployment pipeline support a built-in "simulation stage" type for financial algorithm solutions that runs the algorithm against a market data replay before live activation?	Pre-beta
Q-FIN-004	For tax solutions, the release pipeline must guarantee activation on the legislative effective date — not before, not after. What is the platform mechanism for effective-date-gated activation? How should the system handle the failure scenario where a tax rule set fails its validation gate hours before a mandatory effective date?	Pre-beta

RFC-012 — Agricultural and Agribusiness Solution Domain Extension

Status: Open for Comment **Domain:** Domain Extension

Q-ID	Open Question — Requires Decision	Decision By
Q-AGRI-001	Agricultural solutions include IoT edge deployments (sensors, drones, irrigation controllers) that may operate for months in disconnected environments. What is the OTA update governance model for field-deployed agricultural IoT solutions? How should Qala manage the pipeline for building, testing, and releasing firmware to edge devices with unreliable connectivity?	Pre-beta
Q-AGRI-002	Farm-to-consumer traceability data (crop origin, field coordinates, treatment records, cold chain events) is increasingly required by retailers (Walmart Produce Traceability Initiative) and regulators (EU FIC Regulation). Should Qala provide a native "provenance QR code" export capability that generates a consumer-scannable traceability summary for any agricultural solution batch?	Pre-GA
Q-AGRI-003	Agricultural release cycles are fundamentally calendar-driven (planting season, harvest season, regulatory annual windows). Should Qala's release management support calendar-anchored release scheduling as a first-class feature, where a release is planned relative to a seasonal milestone rather than a continuous deployment cadence?	Pre-beta

RFC-013 — Tools, Toolchains, Libraries, Packages and SDK Domain Extension

Status: Open for Comment **Domain:** Domain Extension

Q-ID	Open Question — Requires Decision	Decision By
Q-TOOL-001	When a library or package publishes a breaking change, all downstream consumers are affected. Should Qala automatically identify all SDEs, pipelines, and solutions that consume a given library version, and generate a migration impact report and optionally draft consumer notification messages?	Pre-GA
Q-TOOL-002	Should Qala's Tool Registry double as a community-discoverable tool marketplace where verified tools	Post-GA

	from any organisation can be published, rated, and adopted? This requires a moderation model, a trust scoring system, and a security verification programme for community-contributed tools.	
Q-TOOL-003	Should Qala generate a full SBOM (Software Bill of Materials) not just for software solutions but for every solution type — listing all tools, libraries, services, data sources, and physical inputs the solution depends on? A "Universal Bill of Materials" could be a significant compliance and supply chain transparency contribution.	Pre-beta

RFC-014 — Service, Consulting and Professional Services Domain Extension

Status: Open for Comment **Domain:** Domain Extension

Q-ID	Open Question — Requires Decision	Decision By
Q-SVC-001	Consulting engagements produce client-confidential deliverables. How should Qala support knowledge capture from confidential engagements for institutional learning — while enforcing client confidentiality? Should Qala support "anonymised knowledge extraction" where engagement learnings are captured with client identifiers removed?	Pre-beta
Q-SVC-002	Service delivery quality metrics (NPS, SLA adherence, outcome achievement) come from external systems (CRM, survey platforms, ticketing systems). Should Qala provide a universal Quality Event Ingestion API that accepts quality metrics from any external system, enabling managed services and consulting firms to benchmark quality without migrating their existing tools?	Pre-alpha build
Q-SVC-003	Should Qala support a "Service Catalogue" solution type — a managed collection of versioned, approved service offerings — as a first-class entity distinct from individual Service solutions? A Service Catalogue has its own governance (catalogue publication approval), versioning (catalogue version), and distribution (who can access the catalogue).	Pre-beta

RFC-015 — Business, Process and Operational Solution Domain Extension

Status: Open for Comment **Domain:** Domain Extension

Q-ID	Open Question — Requires Decision	Decision By
Q-BIZ-001	Business solutions are described in process modelling languages (BPMN, ArchiMate) and natural language documents — not code. Should Qala's SDE and Artefact Registry natively support process model artefact types (BPMN XML, ArchiMate files, Lucidchart exports) with version control and diff comparison, treating them as first-class artefacts equivalent to code?	Pre-alpha build
Q-BIZ-002	Measuring the "quality" of a business solution requires telemetry from the operational environment (process execution time, error rates, adoption metrics) — data that lives in ERP, BPM, and analytics systems outside Qala. Should Qala define a standard operational outcome metrics ingestion model for business	Pre-beta

	solutions, analogous to how software solutions emit metrics to Prometheus?	

RFC-016 — Research, Academic and Scientific Solution Domain Extension

Status: Open for Comment **Domain:** Domain Extension

Q-ID	Open Question — Requires Decision	Decision By
Q-RES-001	Research datasets can be terabytes to petabytes. The Artefact Registry stores metadata and checksums, not the data itself. Should Qala define a "Dataset Pointer" artefact type that stores a verified reference (URL + checksum + access credentials) to externally stored datasets (AWS S3, institutional repositories, Zenodo, CERN), while maintaining version integrity through the platform?	Pre-beta
Q-RES-002	Academic peer review is the quality gate for research publications. Should Qala model peer review as a specialised multi-stage approval workflow in the Release Management system, with stages for submission, editorial desk review, peer assignment, reviewer reports, author response, and final decision?	Pre-beta

RFC-017 — Legal, Tax and Regulatory Solution Domain Extension

Status: Open for Comment **Domain:** Domain Extension

Q-ID	Open Question — Requires Decision	Decision By
Q-LEG-001	Legal and tax solutions change in response to external legislative events — changes initiated outside the organisation. Should Qala integrate with regulatory monitoring services (Thomson Reuters, Bloomberg Law, Wolters Kluwer) to automatically detect relevant legislative changes and generate pre-populated CCRs for the affected solution? This would make regulatory change governance proactive.	Post-GA
Q-LEG-002	Legal solutions often exist in jurisdiction-specific variants of a common base (a contract template with local law governing clauses per jurisdiction). Should Qala's solution model support "localised variants" — instances of a solution that share a common base but differ in jurisdiction-specific attributes — with governed variant creation, change propagation from base to variants, and variant-specific release management?	Pre-beta

13 RFC-018: Universal Compliance Framework Architecture

RFC-018 — Universal Compliance Framework Architecture
Status: Open for Comment **Domain:** Compliance

The Universal Compliance Engine must support dozens of regulatory frameworks across all domains from a single, modular, configurable architecture. The goal is continuous compliance — not periodic audits.

Regulatory Framework	Domain	Automation Potential	Key Evidence Qala Generates
SOC 2 Type II	Software / Cloud	80%+ automatable	Access logs, change records, incident records, availability metrics, backup records.
ISO 27001	All domains	70%+ automatable	Risk assessments, audit logs, policy evidence, asset inventory, change records.
GDPR / CCPA	All with personal data	75%+ automatable	PII classification, consent records, access logs, deletion records, data mapping.
SLSA Level 1–4	Software build	99%+ automatable	Build attestations, dependency digests, environment definitions, signer identity.
FDA 21 CFR Part 11	Pharma / Medical Devices	50% automatable	Electronic signatures, audit trail completeness, data integrity, validation records.
FDA GMP / cGMP	CPG / Pharma / Food	40% automatable	Batch records, equipment calibration, change control records, training records.
MiFID II	Financial Services	55% automatable	Trade audit trail, best execution records, product governance evidence.
Basel III / IV	Banking	45% automatable	Risk model documentation, capital calculation audit, stress test records.
Solvency II	Insurance	45% automatable	Actuarial model governance, capital model documentation, risk audit.
FedRAMP / NIST 800-53	Government / Public Sector	50% automatable	Security control evidence, scan results, configuration records, change management.
REACH / EU CSRD	Manufacturing / CPG	40% automatable	Substance records, supply chain provenance, sustainability data, substance history.
IFRS 9 / GAAP	Financial Reporting	20% automatable	Accounting policy evidence, model documentation, financial audit trail.

Q-ID	Open Question — Requires Decision	Decision By
Q-COMP-001	Should Qala build its own compliance evidence mapping engine (mapping platform events to specific control requirements per framework), or partner with an established GRC platform (ServiceNow GRC, Archer, OneTrust) as the compliance reporting layer? Own engine provides tighter integration and	Pre-GA

	lower cost for customers; GRC partnership leverages existing compliance expertise and reduces build scope.	
Q-COMP-002	For non-software domains, compliance frameworks often require manual evidence that Qala cannot generate automatically (physical laboratory inspections, qualified person signatures, physical plant audits). Should Qala support "hybrid compliance evidence" — combining automatically generated platform evidence with manually uploaded evidence from physical processes — in a single compliance report?	Pre-beta
Q-COMP-003	Should Qala publish "Compliance Packs" as an ecosystem product — versioned, community-contributable compliance evidence mappings for specific frameworks — that can be licensed or contributed by regulatory intelligence firms and compliance specialists?	Post-GA

14 RFC-019 to RFC-022: Integration, Multi-Tenancy, Identity and Strategy

RFC-019 — Physical World Integration: ERP, MES, LIMS and IoT Bridges

Status: Open for Comment **Domain:** Integration

For non-software solution types, Qala manages the information layer. The physical production systems (SAP ERP, Siemens MES, LabVantage LIMS, Honeywell DCS, Plex MES) manage the production layer. Integration is critical and must be an open, governed, auditable API model.

Integration Category	Systems	Integration Pattern	Data Direction
ERP Integration	SAP S/4HANA, Oracle ERP, Microsoft Dynamics	Bidirectional — Qala solution records sync with ERP product records; ERP sends production events to Qala.	Qala -> ERP: specification releases. ERP -> Qala: batch records, production events.
MES Integration	Siemens SIMATIC, Plex, Rockwell	Unidirectional from MES to Qala — production execution events published to Qala event bus.	MES -> Qala: batch start/complete, quality check events, equipment calibration records.
LIMS Integration	LabVantage, STARLIMS, LabWare	Bidirectional — Qala sends test requests; LIMS returns results as quality events.	Qala -> LIMS: test requests on solution release. LIMS -> Qala: test results, CoA.
IoT / Edge Integration	AWS IoT, Azure IoT Hub, Particle	Event stream — device telemetry ingested via IoT platform pre-aggregation; device firmware managed via Qala pipeline.	IoT Platform -> Qala: telemetry, alerts. Qala -> Devices: firmware OTA via pipeline.
Supply Chain / Logistics	SAP TM, Oracle TMS, Manhattan	Unidirectional from logistics to Qala — distribution events published to Qala for provenance tracking.	Logistics -> Qala: shipment events, handoff records, cold chain data.

Q-ID	Open Question — Requires Decision	Decision By
Q-INT-001	Should Qala publish an open "Physical World Integration API" specification that any ERP/MES/LIMS/IoT vendor can implement to emit standardised events to Qala? Or should Qala build and maintain specific point-to-point integration connectors for major systems (SAP, Oracle, Siemens)? Open spec enables ecosystem breadth; point-to-point integrations enable faster enterprise deals.	Pre-GA
Q-INT-002	For IoT integration, agricultural and industrial IoT deployments may generate millions of telemetry events per second. Should Qala's event bus directly accept IoT telemetry, or require an intermediate IoT platform (AWS IoT, Azure IoT Hub) for	Pre-beta

	aggregation, filtering, and pre-processing before events reach the Qala bus?	
Q-INT-003	Physical world integration events (batch records, production events, quality test results) arrive from external systems that Qala does not control. How should these external events be authenticated, provenance-verified, and integrated into the immutable audit trail without compromising the integrity of the platform's own event record?	Pre-alpha build

RFC-020 — Multi-Tenancy, Tenant Isolation and Cross-Tenant Sharing

Status: Open for Comment **Domain:** Platform

Isolation Tier	Model	Regulatory Context	Cost Tier
Community	Shared infrastructure, logical isolation via tenant_id.	Consumer/hobbyist. No material regulation.	Lowest — free or low-cost tier.
Standard	Shared compute, isolated databases (Row-Level Security).	SMB, non-regulated professional use.	Mid — standard subscription.
Professional	Isolated Kubernetes namespace + isolated database.	Mid-market, SOC 2 / ISO 27001 required.	Higher — professional subscription.
Enterprise	Dedicated cluster or VPC. Complete data plane isolation.	Enterprise, HIPAA, MiFID II, government.	Highest — enterprise contract.
Sovereign	Customer-managed cloud or on-premises Qala deployment.	National data sovereignty, air-gap required.	Custom — sovereign deployment contract.

Q-ID	Open Question — Requires Decision	Decision By
Q-MT-001	For supply chain traceability use cases, two separate tenant organisations (a supplier and a buyer) need to share solution traceability records — a supplier's batch records must be visible to the buyer's solution. How should cross-tenant data sharing be implemented with selective visibility, governed consent, and full audit of what was shared, when, and with whom?	Pre-GA
Q-MT-002	Should Qala operate a "public namespace" — a space outside tenant boundaries where open-source libraries, public domain frameworks, and community-contributed resources can be published and consumed by any tenant without cross-tenant trust agreements?	Pre-beta

RFC-021 — Identity, RBAC and Electronic Signature Architecture

Status: Open for Comment **Domain:** Security

Q-ID	Open Question — Requires Decision	Decision By

Q-IAM-001	Non-software personas (CPG Product Developer, Farm Operations Manager, Financial Compliance Officer, Tax Product Manager) have fundamentally different task contexts from software developers. Should the RBAC model include "Persona Profiles" that map roles to simplified, domain-specific portal views — hiding platform capabilities irrelevant to the persona without restricting their actual permissions?	Pre-beta
Q-IAM-002	For regulated domains (pharma under 21 CFR Part 11, financial under MiFID II, legal under eIDAS), electronic signatures with non-repudiation properties are required for specific approval actions. Should Qala support a "Qualified Electronic Signature" integration option using ETSI standards (PAdES, XAdES) or third-party e-signature providers (DocuSign, Adobe Sign) for high-assurance approval workflows?	Pre-GA
Q-IAM-003	Service-to-service authentication within the platform is proposed to use SPIFFE/SPIRE workload identity (mTLS with short-lived X.509 SVIDs). Is this the right choice given SPIFFE's operational complexity relative to shared-secret JWT? What is the minimum viable service identity model for the initial release?	Pre-alpha build

RFC-022 — Open Source, Ecosystem and Commercialisation Strategy

Status: Open for Comment **Domain:** Strategy

Component	Proposed Model	Strategic Rationale
Universal Solution Model Schema	Open Standard (Apache 2.0)	Maximum ecosystem adoption. Enable third-party tooling. Prevent proprietary lock-in.
Platform Kernel	Open Source (Apache 2.0)	Community trust, security auditability, platform stability contributions.
CI/CD Engine	Open Source (Apache 2.0)	Broad developer adoption. Rich integration ecosystem.
Domain Extension Framework	Open Source (Apache 2.0)	Enable community Domain Pack contributions. Third-party vertical extensions.
Domain Packs (core verticals)	Open Source, community-governed	Core domain packs (Software, CPG, Financial, Agricultural, Services) as community assets.
AI Agent	Proprietary + Partner API	Core commercial differentiator. API access for partner integrations at premium tier.
Security Engine (SEM)	Proprietary	Security-critical. Controlled update cadence. Commercial differentiator.
Compliance Reporting Engine	Proprietary + Compliance Pack API	Compliance IP. Partner licensing model for GRC firms.
Enterprise Governance Tier	Proprietary	Advanced RBAC, CCB workflow, policy engine, multi-region — enterprise SKU.

Q-ID	Open Question — Requires Decision	Decision By
Q-OSS-001	Should Qala establish the Universal Solution Model Schema as an open standard governed by an independent foundation (similar to OpenAPI under the OpenAPI Initiative, or AsyncAPI	Pre-GA

	under its own foundation)? Vendor-neutral governance maximises ecosystem trust and adoption but reduces Qala's control over the standard's evolution.	
Q-OSS-002	Should Qala establish a "Domain Pack Marketplace" — a governed registry where third parties (industry consultancies, standards bodies, regulatory agencies, domain experts) can publish, certify, and distribute domain packs? What is the quality assurance, certification, and commercial model for third-party domain packs?	Post-GA
Q-OSS-003	The open source components will attract enterprise deployments that self-host. How should Qala differentiate between self-hosted (open source) and cloud-managed (SaaS) deployments to maintain commercial viability? Should the SaaS tier offer exclusive capabilities unavailable in the open source version, or compete solely on managed service convenience?	Pre-GA

15 Comment Submission and Review Process

Step	Description	Timeline
1. Submit Comment	Submit via issue tracker or document annotation. Reference RFC ID and Question ID in every comment.	Days 1–30
2. Triage	Architecture Review Board categorises: Accepted / Needs Discussion / Declined / Out of Scope.	Days 31–37
3. ARB Sessions	Open RFC discussion sessions for complex or contested questions. Community observers welcome.	Days 38–44
4. Decisions	Every open question receives a documented decision with full rationale. Dissenting views recorded.	Days 45–51
5. RFC Update	Document updated with all decisions. Status per RFC set to Accepted / Modified / Withdrawn.	Days 52–60
6. Design Lock	Accepted RFCs locked. Decisions cascade to SDD v3.0, Domain Pack specifications, and sprint planning.	Day 61+

End of Document — Qala Universal Solution Management OS | RFC v2.0